

Comprehensive Instructions

Author: [Author Name] **Date:** November 4, 2025

Table of Contents

1. API_REFERENCE.md
 2. ENDURANCE_TEST_CONFIG.md
 3. FUTURE_ADDITIONS_ROADMAP.md
 4. HEADLESS_JSON_CLI_ARCHITECTURE.md
 5. IGSOA_2D_ENGINE_IMPLEMENTATION.md
 6. IGSOA_3D_ENGINE_IMPLEMENTATION.md
 7. IGSOA_3D_ROADMAP.md
 8. IGSOA_ANALYSIS_GUIDE.md
 9. IGSOA_COMPLEX_ENGINE_IMPLEMENTATION_PLAN.md
 10. PROJECT_STRUCTURE.md
 11. QIA_IMPLEMENTATION_GUIDE.md
 12. QUICK_START_GUIDE.md
 13. SATP v2 – BLEND_MODE Operational Usage Notes.md
 14. SATP v2 (BLEND_MODE) Reproducibility Standard — 1-page.md
 15. SATP v2 Validation Protocol (external labs).md
 16. SATP_THEORY_PRIMER.md
 17. SATP_v2 (BLEND_MODE) — Canonical Mission (NDJSON).md
 18. SATP_VALIDATION_REPORT.md
 19. SATP_VALIDATION_STATUS.md
 20. SET_IGSOA_STATE_MODES.md
 21. tools_generate_blend_schedule.md
-

API_REFERENCE.md

DASE Engine API Reference

Complete Python & C++ API Documentation

Version: 1.0 **Last Updated:** October 23, 2025 **Performance:** 3.57 ns/op,
280M ops/sec, 4,344x speedup

Table of Contents

1. Python API
 - CPU Features
 - Engine Metrics
 - Analog Universal Node

- Analog Cellular Engine
2. REST API
 3. WebSocket API
 4. C++ API
 5. Examples
-

Python API

Installation

```
# Build the C++ extension
cd src/python
python setup.py build_ext --inplace

# Verify installation
python -c "import dase_engine as de; print('DASE Engine loaded')"
```

Import

```
import dase_engine as de
import numpy as np
```

CPU Features

Check hardware capabilities before using AVX2 features.

CPUFeatures

Static class for CPU capability detection.

Methods

```
has_avx2() → bool Check if CPU supports AVX2 instructions.

if de.CPUFeatures.has_avx2():
    print("AVX2 supported.")
else:
    print("AVX2 not available - performance will be limited")
```

Returns: True if AVX2 is available, False otherwise

`has_fma() → bool` Check if CPU supports FMA (Fused Multiply-Add) instructions.

```
if de.CPUFeatures.has_fma():
    print("FMA supported - extra performance boost!")
```

Returns: True if FMA is available, False otherwise

`print_capabilities() → None` Print all CPU capabilities to console.

```
de.CPUFeatures.print_capabilities()
# Output:
#   CPU Features:
#     AVX2: Supported
#     FMA: Supported
```

Engine Metrics

Performance monitoring and statistics.

EngineMetrics

Class containing performance metrics from C++ engine.

Attributes

Attribute	Type	Description
<code>total_operations</code>	<code>int</code>	Total operations executed
<code>avx2_operations</code>	<code>int</code>	Operations using AVX2 SIMD
<code>node_processes</code>	<code>int</code>	Number of node processing calls
<code>harmonic_generations</code>	<code>int</code>	Harmonic generation count
<code>current_ns_per_op</code>	<code>float</code>	Current nanoseconds per operation
<code>current_ops_per_second</code>	<code>float</code>	Current operations per second
<code>speedup_factor</code>	<code>float</code>	Speedup vs baseline (15,500 ns)
<code>avg_time_ns</code>	<code>float</code>	Average time in nanoseconds
<code>throughput_gflops</code>	<code>float</code>	Throughput in GFLOPS

Methods

`update_performance() → None` Recalculate performance metrics.

```
metrics = engine.get_metrics()
metrics.update_performance()
```

```
print_metrics() → None Print metrics to console in formatted output.  
metrics = engine.get_metrics()  
metrics.print_metrics()  
# Output:  
# D-ASE AVX2 ENGINE METRICS  
# =====  
# Current Performance: 3.57 ns/op  
# Speedup Factor: 4344.48x  
# ...
```

reset() → **None** Reset all metrics to zero.

```
metrics = engine.get_metrics()  
metrics.reset()
```

Analog Universal Node

Single processing node with multiple operations.

AnalogUniversalNodeAVX2

A versatile analog processing node with AVX2 acceleration.

Constructor

```
node = de.AnalogUniversalNodeAVX2()
```

Creates a new node with default state (all zeros).

Attributes

Attribute	Type	Description
x	int16	X coordinate for node positioning
y	int16	Y coordinate for node positioning
z	int16	Z coordinate for node positioning
node_id	int	Unique node identifier
integrator_state	float	Current integrator accumulator value
previous_input	float	Previous input signal value
current_output	float	Current output signal value
feedback_gain	float	Feedback coefficient

Basic Operations

```
amplify(input_signal, gain) → float Multiply input by gain factor.  
output = node.amplify(input_signal=2.0, gain=0.5)  
# output = 1.0
```

Parameters: - `input_signal` (float): Input signal value - `gain` (float): Amplification factor

Returns: Amplified signal ($\text{input} \times \text{gain}$)

Performance: ~0.5 ns per call

```
integrate(input_signal, time_constant) → float Perform time-domain  
integration.  
# Accumulate signal over time  
for i in range(100):  
    output = node.integrate(input_signal=1.0, time_constant=0.1)
```

Parameters: - `input_signal` (float): Input signal to integrate - `time_constant` (float): Integration time constant (seconds)

Returns: Integrated output

Performance: ~1.0 ns per call

Note: Uses thread-local accumulator for state

```
apply_feedback(input_signal, feedback_gain) → float Apply feedback  
from integrator state.  
output = node.apply_feedback(input_signal=1.0, feedback_gain=0.5)
```

Parameters: - `input_signal` (float): Input signal - `feedback_gain` (float): Feedback coefficient

Returns: Input + (`integrator_state` \times `feedback_gain`)

Performance: ~0.5 ns per call

Advanced Operations

```
process_signal_avx2(input_signal, control_signal, aux_signal) →  
float Complete signal processing chain with AVX2 optimization.  
output = node.process_signal_avx2(  
    input_signal=1.0,
```

```
        control_signal=0.5,  
        aux_signal=0.1  
)
```

Parameters: - `input_signal` (float): Primary input signal - `control_signal` (float): Control/modulation signal - `aux_signal` (float): Auxiliary signal

Returns: Processed output signal

Performance: ~3.57 ns per call (with AVX2)

Processing Chain: 1. Amplify input by control 2. Integrate result 3. Blend with auxiliary 4. Apply spectral processing (AVX2) 5. Apply feedback 6. Clamp output to [-10, +10]

Waveform Generation

```
oscillate(frequency_hz, duration_seconds) → List[float] Generate si-  
nusoidal waveform (standard version).
```

```
# Generate 440 Hz tone for 0.1 seconds  
waveform = node.oscillate(frequency_hz=440.0, duration_seconds=0.1)  
# Returns 4,800 samples at 48kHz sample rate
```

Parameters: - `frequency_hz` (float): Frequency in Hertz - `duration_seconds` (float): Duration in seconds

Returns: List of float samples (48kHz sample rate)

Performance: ~0.027 µs per sample (37M samples/sec)

Sample Rate: 48,000 Hz

Range: [-1.0, +1.0] normalized

```
oscillate_np(frequency_hz, duration_seconds) → np.ndarray Generate  
waveform with NumPy zero-copy (more efficient).
```

```
# NumPy zero-copy version - more efficient  
waveform = node.oscillate_np(frequency_hz=440.0, duration_seconds=0.1)  
# Returns NumPy array directly, no Python list conversion
```

Parameters: - `frequency_hz` (float): Frequency in Hertz - `duration_seconds` (float): Duration in seconds

Returns: NumPy array of float32 samples

Performance: 2-3x faster than `oscillate()` (zero-copy)

Benefit: No Python list conversion overhead

Signal Processing

`process_block_frequency_domain(input_block)` → `List[float]` Process signal block using FFT bandpass filter.

`import numpy as np`

Create test signal

`signal = np.sin(np.linspace(0, 2*np.pi, 256))`

Process with FFT bandpass filter

`filtered = node.process_block_frequency_domain(signal.tolist())`

Parameters: - `input_block` (`List[float]`): Input signal samples

Returns: Filtered signal samples

Performance: ~0.07 ms for 256 samples

Filter: Bandpass (keeps middle 50% of spectrum)

Method: FFT → filter → inverse FFT

`process_block_frequency_domain_np(data)` → `np.ndarray` Process signal with NumPy zero-copy (in-place, faster).

`import numpy as np`

Create test signal (float32 array)

`signal = np.sin(np.linspace(0, 2*np.pi, 256)).astype(np.float32)`

Process in-place (modifies signal array)

`filtered = node.process_block_frequency_domain_np(signal)`

filtered and signal are the same object

Parameters: - `data` (`np.ndarray`): NumPy array of float32 samples (modified in-place)

Returns: Same NumPy array (modified)

Performance: Faster than standard version (zero-copy)

Note: Modifies input array in-place!

Batch Processing

```
process_batch(input_signals, control_signals, aux_signals) →  
List[float] Process multiple samples in one call (5-10x faster than loop).
```

```
# Process 1000 samples at once  
inputs = [1.0] * 1000  
controls = [0.5] * 1000  
aux = [0.1] * 1000  
  
outputs = node.process_batch(inputs, controls, aux)  
# Returns list of 1000 processed values
```

Parameters: - `input_signals` (List[float]): Batch of input signals - `control_signals` (List[float]): Batch of control signals - `aux_signals` (List[float]): Batch of auxiliary signals

Returns: List of processed output values

Performance: 5-10x faster than calling `process_signal_avx2()` in a loop

Requirement: All input lists must have the same length

Benefit: Reduces Python function call overhead

State Management

```
set_feedback(feedback_coefficient) → None Set feedback gain coefficient.
```

```
node.set_feedback(0.5) # 50% feedback
```

Parameters: - `feedback_coefficient` (float): Feedback gain (clamped to [-2.0, +2.0])

```
get_output() → float Get current output value.
```

```
output = node.get_output()
```

Returns: Current output signal value

```
get_integrator_state() → float Get current integrator accumulator value.
```

```
state = node.get_integrator_state()
```

Returns: Current integrator state

```
reset_integrator() → None Reset integrator state to zero.  
node.reset_integrator()
```

Analog Cellular Engine

Multi-node engine for large-scale parallel processing.

AnalogCellularEngineAVX2

High-performance engine managing multiple nodes with OpenMP parallelization.

Constructor

```
engine = de.AnalogCellularEngineAVX2(num_nodes=1024)
```

Parameters: - num_nodes (int): Number of processing nodes (default: 1024)

Memory: ~64 bytes per node (64-byte aligned for cache optimization)

Initialization: All nodes initialized with default state

Execution

```
run_mission(cycles) → None Execute processing mission with OpenMP parallelization.
```

```
engine = de.AnalogCellularEngineAVX2(num_nodes=1024)  
engine.run_mission(cycles=3000)  
# Executes 92,160,000 operations (1024 nodes × 30 iterations × 3000 cycles)
```

Parameters: - cycles (int): Number of processing cycles

Performance: 3.57 ns per operation (280M ops/sec)

Parallelization: OpenMP across all CPU cores

Output: Prints metrics to console after completion

```
run_builtin_benchmark(iterations) → None Run built-in benchmark test.
```

```
engine.run_builtin_benchmark(iterations=1000)
```

Parameters: - iterations (int): Number of benchmark iterations

Signal Processing

```
process_signal_wave_avx2(input_signal, control_pattern) → float  
Process single signal with all nodes.
```

```
output = engine.process_signal_wave_avx2(  
    input_signal=1.0,  
    control_pattern=0.5  
)
```

Parameters: - `input_signal` (float): Input signal value - `control_pattern` (float): Control pattern value

Returns: Processed output signal

```
perform_signal_sweep_avx2(frequency) → float Perform frequency sweep  
across nodes.
```

```
output = engine.perform_signal_sweep_avx2(frequency=440.0)
```

Parameters: - `frequency` (float): Sweep frequency in Hz

Returns: Sweep result

```
process_block_frequency_domain(signal_block) → None Process entire  
signal block with FFT (in-place).
```

```
signal = [1.0] * 256  
engine.process_block_frequency_domain(signal) # Modifies signal in-place
```

Parameters: - `signal_block` (List[float]): Signal samples (modified in-place)

Note: Modifies input list in-place!

Utility

```
calculate_inter_node_coupling(node_index) → float Calculate coupling  
factor for specific node.
```

```
coupling = engine.calculate_inter_node_coupling(node_index=10)
```

Parameters: - `node_index` (int): Index of node to analyze

Returns: Coupling factor

```
generate_noise_signal() → float Generate random noise signal.
```

```
noise = engine.generate_noise_signal()
```

Returns: Random noise value

Metrics

```
get_metrics() → EngineMetrics Get performance metrics object.
```

```
metrics = engine.get_metrics()  
print(f"ns/op: {metrics.current_ns_per_op}")  
print(f"Speedup: {metrics.speedup_factor}")
```

Returns: EngineMetrics object with current statistics

```
print_live_metrics() → None Print current metrics to console.
```

```
engine.print_.metrics()
```

```
reset_metrics() → None Reset all metrics to zero.
```

```
engine.reset_metrics()
```

REST API

Flask-based REST API for web integration.

Server Setup

```
cd src/python  
python bridge_server_improved.py  
# Server runs on http://127.0.0.1:5000
```

Endpoints

GET / Serve main web interface.

Response: HTML page

URL: <http://127.0.0.1:5000/>

GET /api/status Get server and engine status.

Response:

```
{  
    "server": "running",  
    "engine_available": true,  
    "timestamp": "2025-10-23T06:50:01",  
    "cpu_capabilities": {  
        "avx2": true,  
        "fma": true  
    },  
    "performance": {  
        "total_requests": 42,  
        "active_simulations": 0,  
        "total_operations": 92160000  
    }  
}
```

GET /api/metrics Get performance metrics.

Response:

```
{  
    "timestamp": "2025-10-23T06:50:01",  
    "server_metrics": {  
        "total_requests": 42,  
        "active_simulations": 0  
    },  
    "engine_available": true  
}
```

POST /api/benchmark Run benchmark test.

Request Body:

```
{  
    "nodes": 1024,  
    "cycles": 1000  
}
```

Response:

```
{  
    "success": true,  
    "benchmark_results": {
```

```

    "nodes": 1024,
    "cycles": 1000,
    "elapsed_seconds": 0.329,
    "cpp_metrics": {
        "total_operations": 30720000,
        "avx2_operations": 30720000,
        "current_ns_per_op": 3.57,
        "current_ops_per_second": 280289072,
        "speedup_factor": 4344.48
    }
}
}

```

POST /api/validate_formula Validate DVSL formula syntax.

Request Body:

```
{
    "formula": "=DELTA(A1, gain=2.0)"
}
```

Response:

```
{
    "valid": true,
    "formula": "=DELTA(A1, gain=2.0)",
    "parsed": {
        "type": "formula",
        "expression": "DELTA(A1, gain=2.0)"
    }
}
```

WebSocket API

Real-time bidirectional communication.

Connection

```
const ws = new WebSocket('ws://127.0.0.1:5000/ws');

ws.onopen = () => {
    console.log('Connected to DASE engine');
};

ws.onmessage = (event) => {
```

```
    const data = JSON.parse(event.data);
    console.log('Received:', data);
};
```

Messages

Client → Server

ping Heartbeat check.

```
ws.send(JSON.stringify({ command: "ping" }));
```

Response: { "type": "pong" }

status Get engine status.

```
ws.send(JSON.stringify({ command: "status" }));
```

Response:

```
{
  "type": "status",
  "engine_available": true,
  "active_simulations": 0
}
```

run Run simulation.

```
ws.send(JSON.stringify({
  command: "run",
  nodes: 1024,
  cycles: 3000
}));
```

Response (start):

```
{
  "type": "run_started",
  "nodes": 1024,
  "cycles": 3000
}
```

Response (done):

```
{
  "type": "run_done",
  "metrics": {
```

```
        "total_operations": 92160000,
        "current_ns_per_op": 3.57,
        "current_ops_per_second": 280289072,
        "speedup_factor": 4344.48
    },
    "timestamp": "2025-10-23T06:50:01"
}
```

Server → Client

hello Initial connection message.

```
{
    "type": "hello",
    "client_id": 140234567890,
    "caps": {
        "avx2": true,
        "fma": true
    },
    "server_time": "2025-10-23T06:50:01"
}
```

error Error notification.

```
{
    "type": "error",
    "error": "Engine not available"
}
```

C++ API

Direct C++ usage (advanced users).

Header Include

```
#include "analog_universal_node_engine_avx2.h"
```

Basic Usage

```
#include "analog_universal_node_engine_avx2.h"
#include <iostream>
```

```

int main() {
    // Check CPU capabilities
    if (!CPUFeatures::hasAVX2()) {
        std::cerr << "AVX2 not available!" << std::endl;
        return 1;
    }

    // Create engine
    AnalogCellularEngineAVX2 engine(1024);

    // Run mission
    engine.runMission(3000);

    // Get metrics
    EngineMetrics metrics = engine.getMetrics();
    std::cout << "ns/op: " << metrics.current_ns_per_op << std::endl;

    return 0;
}

```

Compilation

```

# Compile with AVX2 and OpenMP
g++ -o my_program my_program.cpp \
     analog_universal_node_engine_avx2.cpp \
     -std=c++17 -O2 -mavx2 -fopenmp -lfftw3

# Or with MSVC
cl /EHsc /std:c++17 /O2 /Ob3 /arch:AVX2 /openmp \
    my_program.cpp analog_universal_node_engine_avx2.cpp \
    libfftw3-3.lib

```

Examples

Example 1: Basic Signal Processing

```

import dase_engine as de

# Create node
node = de.AnalogUniversalNodeAVX2()

# Process single signal
output = node.process_signal_avx2(
    input_signal=1.0,
    control_signal=0.5,
)

```

```

        aux_signal=0.1
    )
print(f"Output: {output}")

```

Example 2: Waveform Generation

```

import dase_engine as de
import numpy as np
import matplotlib.pyplot as plt

# Create node
node = de.AnalogUniversalNodeAVX2()

# Generate 440 Hz tone (A4 note)
waveform = node.oscillate_np(frequency_hz=440.0, duration_seconds=0.1)

# Plot
plt.plot(waveform[:1000]) # First 1000 samples
plt.title("440 Hz Sine Wave")
plt.xlabel("Sample")
plt.ylabel("Amplitude")
plt.show()

```

Example 3: Batch Processing

```

import dase_engine as de
import numpy as np

# Create node
node = de.AnalogUniversalNodeAVX2()

# Prepare batch data (1000 samples)
inputs = np.random.randn(1000).tolist()
controls = [0.5] * 1000
aux = [0.1] * 1000

# Process batch (more efficient than a loop)
outputs = node.process_batch(inputs, controls, aux)
print(f"Processed {len(outputs)} samples")

```

Example 4: FFT Filtering

```

import dase_engine as de
import numpy as np

# Create node
node = de.AnalogUniversalNodeAVX2()

```

```

# Create noisy signal
t = np.linspace(0, 1, 1000)
signal = np.sin(2 * np.pi * 10 * t) # 10 Hz sine
noise = np.random.randn(1000) * 0.1
noisy_signal = (signal + noise).astype(np.float32)

# Filter with FFT (in-place)
filtered = node.process_block_frequency_domain_np(noisy_signal)
print(f"Signal filtered (energy reduced by bandpass)")

```

Example 5: Large-Scale Simulation

```

import dase_engine as de
import time

# Create large engine
engine = de.AnalogCellularEngineAVX2(num_nodes=4096)

# Run mission
start = time.time()
engine.run_mission(cycles=5000)
elapsed = time.time() - start

# Get metrics
metrics = engine.get_metrics()
print(f"Elapsed: {elapsed:.2f}s")
print(f"Operations: {metrics.total_operations}")
print(f"Throughput: {metrics.current_ops_per_second/1e6:.1f}M ops/sec")
print(f"Speedup: {metrics.speedup_factor:.1f}x")

```

Example 6: WebSocket Integration

```

// Connect to engine
const ws = new WebSocket('ws://127.0.0.1:5000/ws');

ws.onopen = () => {
    // Run simulation
    ws.send(JSON.stringify({
        command: "run",
        nodes: 1024,
        cycles: 1000
    }));
};

ws.onmessage = (event) => {

```

```

const data = JSON.parse(event.data);

if (data.type === 'run_done') {
    console.log(`Performance: ${data.metrics.current_ns_per_op} ns/op`);
    console.log(`Speedup: ${data.metrics.speedup_factor}x`);
}
};



---



```

Performance Tips

1. Use NumPy Zero-Copy Methods

```

# Slow (Python list conversion)
waveform = node.oscillate(440.0, 0.1)

# Fast (NumPy zero-copy, more efficient)
waveform = node.oscillate_np(440.0, 0.1)

```

2. Use Batch Processing

```

# Slow (Python function call overhead)
outputs = [node.process_signal_avx2(x, c, a) for x, c, a in zip(inputs, controls, aux)]

# Fast (more efficient)
outputs = node.process_batch(inputs, controls, aux)

```

3. Verify AVX2 Support

```

if not de.CPUFeatures.has_avx2():
    print("Warning: AVX2 not available, performance will be limited!")

```

4. Use Appropriate Node Count

```

# Too few: underutilizes CPU cores
engine = de.AnalogCellularEngineAVX2(num_nodes=10)

# Good: matches CPU thread count × workload
engine = de.AnalogCellularEngineAVX2(num_nodes=1024)

# Too many: memory overhead
engine = de.AnalogCellularEngineAVX2(num_nodes=100000)

```

Error Handling

```
import dase_engine as de

try:
    # Create node
    node = de.AnalogUniversalNodeAVX2()

    # Process signal
    output = node.process_signal_avx2(1.0, 0.5, 0.1)

except ImportError as e:
    print(f"Engine not available: {e}")
    print("Build with: python setup.py build_ext --inplace")

except RuntimeError as e:
    print(f"Runtime error: {e}")

except Exception as e:
    print(f"Unexpected error: {e}")
```

Version Information

Current Version: 1.0

Performance Metrics: - Operation Time: 3.57 ns - Throughput: 280M ops/sec - Speedup: 4,344x vs baseline - AVX2 Usage: 100%

Compiler Flags: - /O2 - Optimization level 2 - /Ob3 - Aggressive inlining - /Oi - Intrinsic functions - /Ot - Favor speed - /arch:AVX2 - AVX2 instructions - /fp:fast - Fast floating-point - /GL - Whole program optimization - /LTCG - Link-time code generation

Support

Issues: Report at project repository **Documentation:** See docs/ directory
Examples: See tests/ directory

End of API Reference

...and so on for all the other files.